

# EE-147 Lab 3

## Histogram

Cody Simons  
861177050

May 10, 2018

## 1 Introduction

In this lab we learn how to use the atomicAdd operation. Our use case is calculating a histogram for a random vector. It is important that we use the atomicAdd operation rather than normal add to increment, because it ensures that no other may access the values in the bins while another process is working with them. This prevents a lot of undefined behaviors that would be prevalent if memory read-modify-writes were randomly interleaved.

## 2 Questions

1. Use visual profiler to report relevant statistics (e.g. utilization and memory hierarchy related) about the execution of your kernels. Did you find any surprising results?

I was unable to work with the visual profiler, but using the command line version I was able to find some interesting results. There are no global store request which is odd, because I would expect that when all the private histograms are added together that would generate global store requests.

==13880== Profiling application: ./histogram							
==13880== Profiling result:							
	Type	Time(%)	Time	Calls	Avg	Min	
	Max	Name					
GPU activities:	69.12%	1.6636ms	1	1.6636ms	1.6636ms	1.6636	
ms [CUDA memcpy HtoD]	30.51%	734.34us	1	734.34us	734.34us	734.34	
		us histogram_kernel(unsigned int*, unsigned int*, unsigned int, unsigned int)					
	0.19%	4.5440us	1	4.5440us	4.5440us	4.5440	
	us [CUDA memset]						
	0.19%	4.5120us	1	4.5120us	4.5120us	4.5120	
	us [CUDA memcpy DtoH]						
API calls:	98.23%	211.18ms	2	105.59ms	130.71us	211.05	
ms cudaMalloc	0.86%	1.8527ms	2	926.33us	40.074us	1.8126	
ms cudaMemcpy	0.41%	879.03us	4	219.76us	3.7380us	743.57	
	us cudaDeviceSynchronize						
	0.33%	703.52us	2	351.76us	119.70us	583.82	
	us cudaFree						
	0.07%	148.34us	94	1.5780us	187ns	57.658	
	us cuDeviceGetAttribute						
	0.06%	136.36us	1	136.36us	136.36us	136.36	
	us cuDeviceTotalMem						

0.02%	45.102 us	1	45.102 us	45.102 us	45.102
us	cudaLaunch				
0.01%	21.961 us	1	21.961 us	21.961 us	21.961
us	cuDeviceGetName				
0.01%	15.589 us	1	15.589 us	15.589 us	15.589
us	cudaMemset				
0.00%	6.3150 us	4	1.5780 us	179 ns	5.3240
us	cudaSetupArgument				
0.00%	2.8810 us	2	1.4400 us	318 ns	2.5630
us	cuDeviceGet				
0.00%	2.8700 us	3	956 ns	200 ns	1.9420
us	cuDeviceGetCount				
0.00%	887 ns	1	887 ns	887 ns	887
ns	cudaConfigureCall				

2. What, if any, limitations are there on m and n for your implementation? Explain these limitations and how you may overcome them with a different implementation.

The lower bound would be that you need at least one bin and one element in your input. This can't really be changed due the nature of the problem. The upper bound would be limited by two different factors. The number of bins would have to be less than the maximum amount of shared memory divided by the size of an unsigned int. The number of elements in the input vector has a similar constraint, but based on the size of global memory. Unfortunately you can't change the bin size limit because atomicAdd only supports 32 bit values. You could increase the number of elements in the input by using a smaller data size.

## 3 Code

Below is my kernel for generating the histogram. Each SM has its own private histogram that is shared between all the thread in the warp. These are all then summed together to get the final result.

### 3.1 Kernel.cu

```

1  /*****
2  *CR
3  *cr      (C) Copyright 2010 The Board of Trustees of the
4  *cr      University of Illinois
5  *cr      All Rights Reserved
6  *CR
7  *****/
8
9  // Define your kernels in this file you may use more than one kernel if you
10 // need to
11
12 // INSERT KERNEL(S) HERE
13 #include <stdio.h>
14
15 __global__ void histogram_kernel(unsigned int *inputs, unsigned int *bins, unsigned int num_elements, unsigned int *sum) {
16     extern __shared__ unsigned int private_histogram[];
17
18     int bx = blockIdx.x; int bd = blockDim.x;
19     int tx = threadIdx.x; int gd = gridDim.x;
20
21     int i=0;
22     while(i*bd+tx < num_bins) {
23         private_histogram[i*bd+tx] = 0;
24         i++;
25     }
26     __syncthreads();
27
28     int index = bx*bd+tx;
29     int stride = bd*gd;

```

```

30     i=0;
31     while(i*stride+index < num_elements) {
32         atomicAdd(private_histogram+inputs[i*stride+index],1);
33         i++;
34     }
35     __syncthreads();
36
37     i=0;
38     while(i*bd+tx < num_bins) {
39         atomicAdd(bins+i*bd+tx,private_histogram[i*bd+tx]);
40         i++;
41     }
42
43     return;
44 }
45
46 /*****
47 Setup and invoke your kernel(s) in this function. You may also allocate more
48 GPU memory if you need to
49 *****/
50 void histogram(unsigned int* input, unsigned int* bins, unsigned int num_elements,
51               unsigned int num_bins) {
52
53     // INSERT CODE HERE
54     dim3 gridDim(30,1,1);
55     dim3 blockDim(32,1,1);
56
57     histogram_kernel<<<gridDim, blockDim, num_bins*sizeof(unsigned int)>>>(input, bins, num_elements, num_bins);
58 }

```

## 4 Program Output

```

bender /home/eemaj/csimons/EE-147/Lab 3 $ ./histogram

Setting up the problem...0.014242 s
    Input size = 1000000
    Number of bins = 4096
Allocating device variables...0.319838 s
Copying data from host to device...0.001842 s
Launching kernel...0.000785 s
Copying data from device to host...0.000043 s
Verifying results...TEST PASSED

```