

EE-147 Lab 4

CUDA Streams

Cody Simons
861177050

May 30, 2018

1 Introduction

In this lab we revisit our vector addition code from lab one. The outcome should be the same, but this time we are modifying our code to make use of streaming. This allows us to begin calculations before all of the data is written to the GPU, which helps to offset the cost of moving data to and from the GPU.

2 Results

Comparing the streaming and the non-streaming version of vector add, we can see that the time taken to do the computation is a lot lower. This is because the kernel doesn't have to wait for the memory copy to complete before it begins its computation. There is some overhead to create the streams, so the streaming version does take longer to run. A comparison of the computation time and memcpy time is shown in fig 1. The time line generated by the Nvidia visual profiler is also shown in fig 2 and in fig 3. We can see that in the non-streaming version, the computation and memory copies are overlapping and take up much less space.

3 Questions

1. What is the speed up between the non-Stream and Stream version of Vector Add? Where do the improvement comes from?
The speed up comes from the stream essentially pipelining the vector add kernel. It allows the vector addition to start while the data is still streaming to the GPU. Then before the kernel finishes running it can move the result back to host.
2. How can data transfers be further optimized?
You could initialize the host variable in pinned memory so that the you don't have to copy the variable into pinned memory before you start the transfer.
3. Do ordering of various CUDA API calls on the host side matter when implementing streams? Why or why not?

non-stream vector add	116.21 ms
memcpy HtoD	1.74 ms
memcpy DtoH	1.92 ms
kernel	0.103 ms
stream vector add	181.69 ms
memcpy HtoD	1.51 ms
memcpy DtoH	0.65 ms
kernel	0.08 ms

Figure 1: Comparison of streaming and non-streaming memory

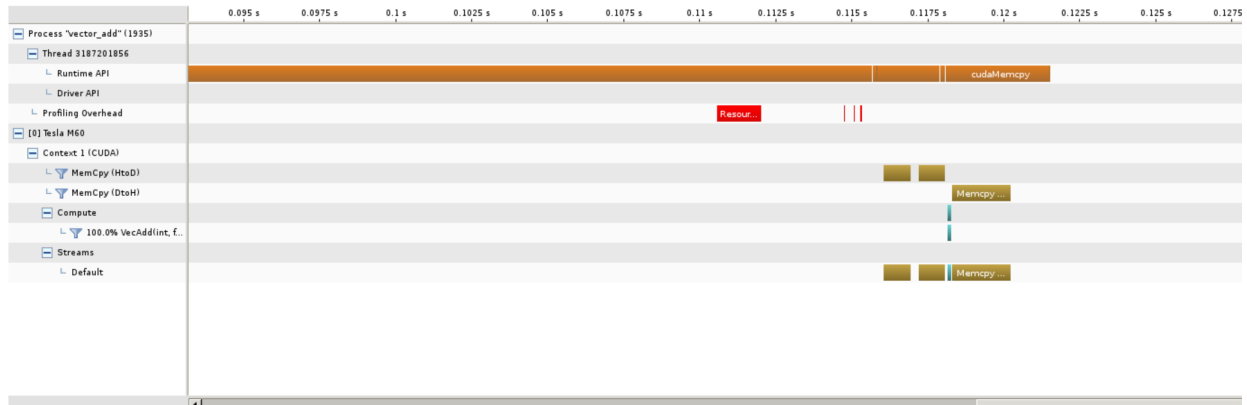


Figure 2: Non-streaming Vector Add

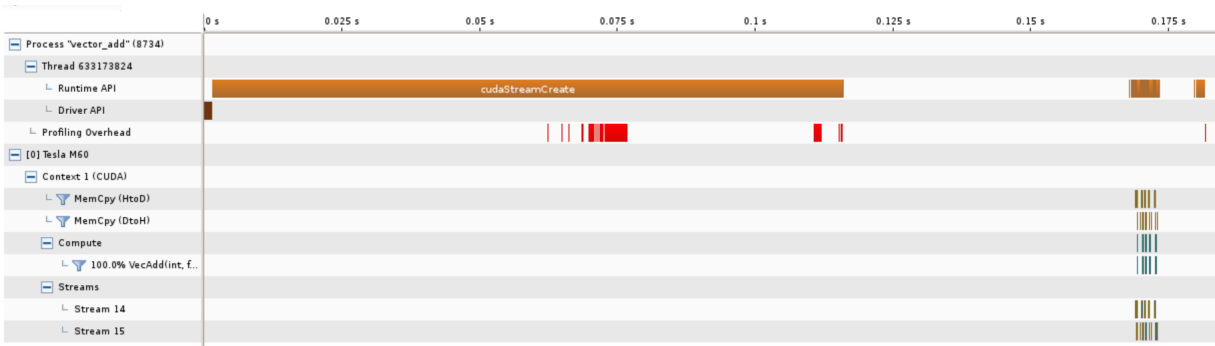


Figure 3: Streaming Vector Add

Yes, generally you need to start streaming the data before you call the kernel. It would also be better to start streaming the data for both stream before you start the kernels, so that you maximize how much is happening in parallel.

4 Code

In this lab the kernel code is reused from lab 1 with no modification, so it is not shown. The main file is modified to create two streams and split the computation between the two streams. The code is shown below.

4.1 main.cu

```

1  /*****
2  *CR
3  *CR      (C) Copyright 2010 The Board of Trustees of the
4  *CR      University of Illinois
5  *CR      All Rights Reserved
6  *CR
7  *****/
8
9  #include <stdio.h>
10 #include <stdlib.h>
11 #include "kernel.cu"
12 #include "support.cu"
13
14 #define SEG_SIZE 100000
15 #define BLOCK_SIZE 256
16
17 int main (int argc, char *argv[])
18 {
19     unsigned VecSize;

```

```

20     if (argc == 1) {
21         VecSize = 1000000;
22     } else if (argc == 2) {
23         VecSize = atoi(argv[1]);
24     } else {
25         printf("\nOh no!\nUsage: ./vecAdd <Size>");
26         exit(0);
27     }
28
29     //set standard seed
30     srand(217); //Default value 217, DO NOT TOUCH
31
32     Timer timer;
33     cudaError_t cuda_ret;
34
35     // Initialize host variables -----
36
37     printf("\nSetting up the problem...\n"); fflush(stdout);
38     startTime(&timer);
39
40     cudaStream_t stream0, stream1;
41     printf("Creating stream0...\n"); fflush(stdout);
42     cuda_ret = cudaStreamCreate(&stream0);
43     if(cuda_ret != cudaSuccess) {
44         printf("Failed to create stream 0, exiting"); fflush(stdout);
45         return 0;
46     }
47     printf("Stream 0 created!\nCreating stream1...\n"); fflush(stdout);
48     cuda_ret = cudaStreamCreate(&stream1);
49     if(cuda_ret != cudaSuccess) {
50         printf("Failed to create stream 0, exiting"); fflush(stdout);
51         return 0;
52     }
53     printf("Stream 1 created!\n"); fflush(stdout);
54
55     float *A_d0, *B_d0, *C_d0; // device memory for stream 0
56     float *A_d1, *B_d1, *C_d1; // device memory for stream 1
57
58     float *A_h, *B_h, *C_h;
59     size_t d0_sz, d1_sz;
60
61     dim3 dim_grid, dim_block;
62
63     d0_sz = VecSize/2;
64     d1_sz = (VecSize-1)/2 + 1;
65
66     A_h = (float*) malloc(sizeof(float)*VecSize);
67     for (unsigned int i=0; i < VecSize; i++) { A_h[i] = (rand()%100)/100.00; }
68
69     B_h = (float*) malloc(sizeof(float)*VecSize);
70     for (unsigned int i=0; i < VecSize; i++) { B_h[i] = (rand()%100)/100.00; }
71
72     C_h = (float*) malloc(sizeof(float)*VecSize);
73
74     stopTime(&timer); printf("%f s\n", elapsedTime(timer));
75     printf("      size Of vector: %u x %u\n ", VecSize,1);
76
77     // Allocate device variables -----
78
79     printf("Allocating device variables..."); fflush(stdout);
80     startTime(&timer);
81
82     //INSERT CODE HERE
83     cudaMalloc((void **) &A_d0, sizeof(float)*d0_sz);
84     cudaMalloc((void **) &B_d0, sizeof(float)*d0_sz);
85     cudaMalloc((void **) &C_d0, sizeof(float)*d0_sz);
86
87     cudaMalloc((void **) &A_d1, sizeof(float)*d1_sz);
88     cudaMalloc((void **) &B_d1, sizeof(float)*d1_sz);
89     cudaMalloc((void **) &C_d1, sizeof(float)*d1_sz);
90

```

```

91 cudaDeviceSynchronize();
92 stopTime(&timer); printf("%f s\n", elapsedTime(timer));
93
94 // Run streams here -----
95
96 printf("Running streaming operations..."); fflush(stdout);
97 startTime(&timer);
98
99 dim3 DimGrid((SEG_SIZE-1)/BLOCK_SIZE+1,1,1);
100 dim3 DimBlock(BLOCK_SIZE,1,1);
101
102 for(int i=0;i<VecSize;i+=SEG_SIZE*2) {
103     if(i+SEG_SIZE*2<VecSize) {
104         cudaMemcpyAsync(A_d0, A_h+i, SEG_SIZE*sizeof(float),
105                         cudaMemcpyHostToDevice, stream0);
106         cudaMemcpyAsync(B_d0, B_h+i, SEG_SIZE*sizeof(float),
107                         cudaMemcpyHostToDevice, stream0);
108         cudaMemcpyAsync(A_d1, A_h+i+SEG_SIZE, SEG_SIZE*sizeof(float),
109                         cudaMemcpyHostToDevice, stream1);
110         cudaMemcpyAsync(B_d1, B_h+i+SEG_SIZE, SEG_SIZE*sizeof(float),
111                         cudaMemcpyHostToDevice, stream1);
112     } else {
113         cudaMemcpyAsync(A_d0, A_h+i, (VecSize-i)/2*sizeof(float),
114                         cudaMemcpyHostToDevice, stream0);
115         cudaMemcpyAsync(B_d0, B_h+i, (VecSize-i)/2*sizeof(float),
116                         cudaMemcpyHostToDevice, stream0);
117         cudaMemcpyAsync(A_d1, A_h+i+(VecSize-i)/2, ((VecSize-i-1)/2+1)*sizeof(float),
118                         cudaMemcpyHostToDevice, stream1);
119         cudaMemcpyAsync(B_d1, B_h+i+(VecSize-i)/2, ((VecSize-i-1)/2+1)*sizeof(float),
120                         cudaMemcpyHostToDevice, stream1);
121     }
122
123     VecAdd<<<DimGrid, DimBlock, 0, stream0>>>(d0_sz, A_d0, B_d0, C_d0);
124     VecAdd<<<DimGrid, DimBlock, 0, stream1>>>(d1_sz, A_d1, B_d1, C_d1);
125
126     if(i+SEG_SIZE*2<VecSize) {
127         cudaMemcpyAsync(C_h+i, C_d0, SEG_SIZE*sizeof(float),
128                         cudaMemcpyDeviceToHost, stream0);
129         cudaMemcpyAsync(C_h+i+SEG_SIZE, C_d1, SEG_SIZE*sizeof(float),
130                         cudaMemcpyDeviceToHost, stream1);
131     } else {
132         cudaMemcpyAsync(C_h+i, C_d0, (VecSize-i)/2*sizeof(float),
133                         cudaMemcpyDeviceToHost, stream0);
134         cudaMemcpyAsync(C_h+i+(VecSize-i)/2, C_d1, ((VecSize-i-1)/2+1)*sizeof(float),
135                         cudaMemcpyDeviceToHost, stream1);
136     }
137 }
138
139 cudaDeviceSynchronize();
140 stopTime(&timer); printf("%f s\n", elapsedTime(timer));
141
142 // Verify correctness -----
143
144 printf("Verifying results..."); fflush(stdout);
145
146 verify(A_h, B_h, C_h, VecSize);
147
148
149 // Free memory -----
150
151 free(A_h);
152 free(B_h);
153 free(C_h);
154
155 //INSERT CODE HERE
156 cudaFree(A_d0);
157 cudaFree(B_d0);
158 cudaFree(C_d0);
159
160 cudaFree(A_d1);
161 cudaFree(B_d1);

```

```
162     cudaFree(C_d1);
163
164     return 0;
165
166 }
```

5 Program Output

The output is shown below, the verification output is suppressed, because it would print literally a million elements.

```
bender /home/eemaj/csimons/EE-147/Lab 4 $ ./vector-add
```

```
Setting up the problem...
Creating stream0...
Stream 0 created!
Creating stream1...
Stream 1 created!
0.498974 s
    size Of vector: 1000000 x 1
    Allocating device variables...0.000967 s
Running streaming operations...0.005558 s
Verifying results...
TEST PASSED
```